

CareerMap AI

A decision-support platform that turns career confusion into a clear, personalized path.

3 sec to first insight	7 clusters career coverage	+155% user clarity lift	0 hallucinations
--	--	---	--

Role	Product Manager (end-to-end ownership)
Timeline	6 weeks · Discovery → MVP → Validation
Team	1 PM · 1 Designer · 2 Engineers
Surface	Web app · React + TypeScript
Status	Live MVP · Validating with early users

"Most resume tools give you a score. CareerMap gives you a **decision**."

Students don't lack tools — they lack clarity.

We talked to 23 final-year students and early-career switchers. The pattern was sharper than expected: they were drowning in resume scanners, course platforms, and LinkedIn advice — yet none of these tools answered the one question that actually mattered: **“What should I do next?”**

What users actually said

"My resume score is 78. Cool. Now what?"

— B.Tech, 4th year

"I have 40 skills listed. I don't know which 5 actually matter for a Product role."

— Career switcher, marketing → PM

"Every course platform tells me to learn everything. I have 3 months."

— M.Sc graduate

The transparency gap

Existing tools	What's missing	User cost
Resume scanners	Why this score?	No trust
Job boards	Am I ready?	Wasted applications
Course platforms	Which skill first?	3-month detours
LinkedIn advice	Personalised to me?	Generic noise

Who we're building for.

Aspiring Aman Final-year engineering student

Knows he wants 'tech' but not which role. Has projects, no clear narrative.

JTBD: "Tell me which role I'm closest to — and the 3 things to fix this month."

Switching Sara 2 yrs marketing → wants Product

Has transferable skills but no PM signal on her resume.

JTBD: "Show me the realistic path, not the dream path."

Returning Riya Career break, re-entering tech

Confident in fundamentals, anxious about 'gap' perception.

JTBD: "Help me prove I'm still ready — fast."

Goals & success metrics

Goal	Metric	Target
Reduce career confusion	% users selecting a target role within 1 session	≥ 70%
Replace generic scores	% users who can explain their score back	≥ 80%
Drive concrete next action	Roadmap tasks started within 7 days	≥ 3 tasks/user
Build trust	Self-reported clarity (1–5)	≥ 4.2

From score to **decision**.

CareerMap AI is structured around one principle: every screen must reduce a decision, not just display data. The product flows in four deliberate steps — Upload, Map, Compare, Act.

1. Upload	Resume in. Skills, projects, experience parsed.
2. Map	Pick from 7 career clusters → role-specific scoring.
3. Compare	Side-by-side roles. Reality-check vs aspiration.
4. Act	Personalised roadmap with 30/60/90-day tasks.

Design principles

- **Show the math.** Every score breaks down into Skills (40%) · Projects (30%) · Experience (30%).
- **One decision per screen.** No dashboards-of-everything.
- **Deterministic > clever.** Rules-based engine first, LLM only where it can't hallucinate.
- **Aspiration meets reality.** Always pair the 'dream role' with a 'best-fit role'.

Transparent by design.

We rejected the 'magic black-box score' pattern. The Readiness Score is a weighted composite users can mentally reproduce.

Readiness Score formula

Score = (Skills match × **0.40**) + (Projects relevance × **0.30**) + (Experience fit × **0.30**)

Each sub-score is 0–100 and visible to the user with the exact contributing items.

Weight rationale

Component	Weight	Why this weight
Skills match	40%	Strongest predictor of role-fit per user interviews + recruiter signal.
Projects relevance	30%	Proves applied ability; differentiates similar resumes.
Experience fit	30%	Captures domain context — but capped to avoid penalising switchers.

Edge cases we explicitly handled

- **Empty resume:** Show diagnostic flow instead of a 0/100 that demoralises.
- **Over-qualified:** Cap experience contribution; surface 'consider senior roles' nudge.
- **Skill spam (40+ listed):** Weight only skills with project evidence.
- **Career switchers:** Re-weight transferable skills using cross-cluster mapping.

Choosing boring tech on purpose.

Rules-based Advisor (vs LLM)

Decision: Picked deterministic recommendations over LLM generation.

Impact: Zero hallucinations. Users trust the 'why' explanation. LLM reserved for v2 conversational layer.

Curated role taxonomy (vs scraping live job posts)

Decision: Hand-built 7 clusters × ~6 roles with explicit skill weights.

Impact: Slower to expand, but every recommendation is explainable. Live API in v2.

Client-side scoring (vs backend)

Decision: All scoring runs in the browser on parsed JSON.

Impact: Instant feedback, zero infra cost for MVP. Privacy bonus: resume never leaves device.

Compare 2 roles only (vs N)

Decision: Forced binary comparison.

Impact: Reduces analysis paralysis — the #1 problem from user research.

Four screens, one decision each.

Upload → Parse

Drag-drop resume. Parser extracts skills, projects, experience. User confirms in 5 seconds.

Decision: **Is my data correct?**

Select Role → Cluster Picker

7 clusters as cards. Each cluster expands to roles with skill-weight previews.

Decision: **Which role do I want to target?**

Dashboard → Readiness Breakdown

Circular score + 3 sub-scores + 'why this score' panel + closest matching role.

Decision: **How ready am I, really?**

Advisor → Compare & Roadmap

Side-by-side: Aspiration role vs Best-fit role. 30/60/90-day roadmap auto-generated.

Decision: **What do I do this week?**

Did it actually move the needle?

We ran 14 moderated tests after MVP launch. Same 3 tasks, before-and-after CareerMap, self-reported clarity on a 1–5 scale.

1.8 → 4.6 Clarity score (1–5)	+155% Lift in 'I know my next step'	11/14 Selected a target role in 1 session	4.7 Trust score on 'why this'
---	---	---	---

North-star & guardrails

Type	Metric	Why
North star	Roadmap tasks completed in first 14 days	Proxy for real-world action, not vanity engagement.
Input	% users completing role selection	Funnel health.
Quality	Self-reported clarity (post-session)	We optimise for confidence, not time-on-page.
Guardrail	% users who skip score explanation	If high → we lost transparency. Currently 8%.

"This is the first time a tool actually told me what to **do**, not just what I lack."

— User #9, post-test interview

Where v2 goes.

Horizon	Initiative	Why now
Now (Q1)	Real-time NLP resume parsing	Cut confirmation step; reduce parse errors below 3%.
Next (Q2)	Live job-board API (LinkedIn / Naukri)	Convert readiness into actual applications.
Next (Q2)	Peer benchmarking ('users like you')	Social proof drives roadmap adherence.
Later (Q3)	Conversational Advisor (LLM, grounded)	Once rules cover 80% — LLM only fills gaps.
Later (Q3)	Recruiter-side preview ('resume X-ray')	Two-sided validation; new growth surface.

What I'd do differently

- Run the cluster taxonomy past 3 recruiters **before** building UI — would've saved 1 week.
- Ship the Compare flow in week 2, not week 5 — it became the most-loved feature.
- Instrument 'explain the score back to me' as a true funnel step, not a survey question.

What this project taught me as a PM.

Transparency is a feature, not a footnote.

The single highest-impact decision was making the scoring formula visible. Every other piece of trust hung off it.

Decision-support beats data-display.

Users don't want more charts. They want fewer, sharper choices. We removed 4 dashboard widgets and clarity went up.

Pick the boring tech that lets you sleep.

Deterministic rules over LLMs for v1 meant zero late-night hallucination bug reports — and a product I could actually defend in user interviews.

Edge cases are the product.

Empty resumes, switchers, returners — handling these well is what separated us from 'just another scanner.'

Thank you for reading.

CareerMap AI · Built with intent, shipped with evidence.